

Practice 2: Hyperparameter Tuning of Initial Practice (Parkinson's Disease)

Statistical Learning with Deep Artificial Neural Networks

Alexander J Ohrt

11 mars, 2022

Contents

1	Introduction	2
2	Load the Parkinsons Data	2
3	Description of the Variables	3
4	Create the Binary Variable of Parkinson's Severity	4
5	Normalization	4
6	Separation into Train and Test Data	5
7	Implementation of a Dense DNN	5
7.1	Using One Output Node for Classification	5
8	Add Callbacks	7
9	Predictions	9

1 Introduction

We are working on a dataset describing Parkinson's disease. Click [here](#) for more information regarding the dataset. In short, the dataset is composed of a range of biomedical voice measurements from 42 people with early-stage Parkinson's disease recruited to a six-month trial of a telemonitoring device for remote symptom progression monitoring.

The main objective is to predict the severity of Parkinson's disease based on the data. More details are given in the following.

2 Load the Parkinsons Data

```
data <- read.csv("../practice_1/parkinsons_updrs.data")
str(data)
```

```
#> 'data.frame':    5875 obs. of  22 variables:
#> $ subject.      : int  1 1 1 1 1 1 1 1 1 1 ...
#> $ age           : int  72 72 72 72 72 72 72 72 72 72 ...
#> $ sex           : int  0 0 0 0 0 0 0 0 0 0 ...
#> $ test_time     : num  5.64 12.67 19.68 25.65 33.64 ...
#> $ motor_UPDRS   : num  28.2 28.4 28.7 28.9 29.2 ...
#> $ total_UPDRS   : num  34.4 34.9 35.4 35.8 36.4 ...
#> $ Jitter...     : num  0.00662 0.003 0.00481 0.00528 0.00335 0.00353 0.00422 0.00476 0.00432 0.00496 ...
#> $ Jitter.Abs.   : num  3.38e-05 1.68e-05 2.46e-05 2.66e-05 2.01e-05 ...
#> $ Jitter.RAP    : num  0.00401 0.00132 0.00205 0.00191 0.00093 0.00119 0.00212 0.00226 0.00156 0.002 ...
#> $ Jitter.PPQ5   : num  0.00317 0.0015 0.00208 0.00264 0.0013 0.00159 0.00221 0.00259 0.00207 0.00253 ...
#> $ Jitter.DDP    : num  0.01204 0.00395 0.00616 0.00573 0.00278 ...
#> $ Shimmer       : num  0.0256 0.0202 0.0168 0.0231 0.017 ...
#> $ Shimmer.dB.   : num  0.23 0.179 0.181 0.327 0.176 0.214 0.445 0.212 0.371 0.31 ...
#> $ Shimmer.APQ3  : num  0.01438 0.00994 0.00734 0.01106 0.00679 ...
#> $ Shimmer.APQ5  : num  0.01309 0.01072 0.00844 0.01265 0.00929 ...
#> $ Shimmer.APQ11: num  0.0166 0.0169 0.0146 0.0196 0.0182 ...
#> $ Shimmer.DDA   : num  0.0431 0.0298 0.022 0.0332 0.0204 ...
#> $ NHR           : num  0.0143 0.0111 0.0202 0.0278 0.0116 ...
#> $ HNR           : num  21.6 27.2 23 24.4 26.1 ...
#> $ RPDE          : num  0.419 0.435 0.462 0.487 0.472 ...
#> $ DFA           : num  0.548 0.565 0.544 0.578 0.561 ...
#> $ PPE           : num  0.16 0.108 0.21 0.333 0.194 ...
```

```
summary(data)
```

```
#>      subject.      age      sex      test_time
#> Min.   : 1.00   Min.   :36.0   Min.   :0.0000   Min.   : -4.263
#> 1st Qu.:10.00   1st Qu.:58.0   1st Qu.:0.0000   1st Qu.: 46.847
#> Median :22.00   Median :65.0   Median :0.0000   Median : 91.523
#> Mean   :21.49   Mean   :64.8   Mean   :0.3178   Mean   : 92.864
#> 3rd Qu.:33.00   3rd Qu.:72.0   3rd Qu.:1.0000   3rd Qu.:138.445
#> Max.   :42.00   Max.   :85.0   Max.   :1.0000   Max.   :215.490
#>      motor_UPDRS      total_UPDRS      Jitter...      Jitter.Abs.
#> Min.   : 5.038   Min.   : 7.00   Min.   :0.000830   Min.   :2.250e-06
#> 1st Qu.:15.000   1st Qu.:21.37   1st Qu.:0.003580   1st Qu.:2.244e-05
#> Median :20.871   Median :27.58   Median :0.004900   Median :3.453e-05
#> Mean   :21.296   Mean   :29.02   Mean   :0.006154   Mean   :4.403e-05
#> 3rd Qu.:27.596   3rd Qu.:36.40   3rd Qu.:0.006800   3rd Qu.:5.333e-05
#> Max.   :39.511   Max.   :54.99   Max.   :0.099990   Max.   :4.456e-04
```

```

#>      Jitter.RAP      Jitter.PPQ5      Jitter.DDP      Shimmer
#> Min.   :0.000330   Min.   :0.000430   Min.   :0.000980   Min.   :0.00306
#> 1st Qu.:0.001580   1st Qu.:0.001820   1st Qu.:0.004730   1st Qu.:0.01912
#> Median :0.002250   Median :0.002490   Median :0.006750   Median :0.02751
#> Mean   :0.002987   Mean   :0.003277   Mean   :0.008962   Mean   :0.03404
#> 3rd Qu.:0.003290   3rd Qu.:0.003460   3rd Qu.:0.009870   3rd Qu.:0.03975
#> Max.   :0.057540   Max.   :0.069560   Max.   :0.172630   Max.   :0.26863
#>      Shimmer.dB.      Shimmer.APQ3      Shimmer.APQ5      Shimmer.APQ11
#> Min.   :0.026       Min.   :0.00161       Min.   :0.00194       Min.   :0.00249
#> 1st Qu.:0.175       1st Qu.:0.00928       1st Qu.:0.01079       1st Qu.:0.01566
#> Median :0.253       Median :0.01370       Median :0.01594       Median :0.02271
#> Mean   :0.311       Mean   :0.01716       Mean   :0.02014       Mean   :0.02748
#> 3rd Qu.:0.365       3rd Qu.:0.02057       3rd Qu.:0.02375       3rd Qu.:0.03272
#> Max.   :2.107       Max.   :0.16267       Max.   :0.16702       Max.   :0.27546
#>      Shimmer.DDA      NHR      HNR      RPDE
#> Min.   :0.00484       Min.   :0.000286       Min.   : 1.659       Min.   :0.1510
#> 1st Qu.:0.02783       1st Qu.:0.010955       1st Qu.:19.406       1st Qu.:0.4698
#> Median :0.04111       Median :0.018448       Median :21.920       Median :0.5423
#> Mean   :0.05147       Mean   :0.032120       Mean   :21.680       Mean   :0.5415
#> 3rd Qu.:0.06173       3rd Qu.:0.031463       3rd Qu.:24.444       3rd Qu.:0.6140
#> Max.   :0.48802       Max.   :0.748260       Max.   :37.875       Max.   :0.9661
#>      DFA      PPE
#> Min.   :0.5140       Min.   :0.02198
#> 1st Qu.:0.5962       1st Qu.:0.15634
#> Median :0.6436       Median :0.20550
#> Mean   :0.6532       Mean   :0.21959
#> 3rd Qu.:0.7113       3rd Qu.:0.26449
#> Max.   :0.8656       Max.   :0.73173

```

3 Description of the Variables

As we have seen above, the dataset contains 5875 rows, i.e. 5875 measurements. The columns consist of **patient ID**, **age**, **sex**, **time interval since enrollment date**, **motor_UPDRS**, **total_UPDRS** and 16 voice biomedical measurements. The variables are the following

- subject. - The patient ID. Integer that uniquely identifies each subject.
- age - Age of each subject.
- sex - Gender of the subject; '0' = male and '1' = female.
- test_time - Time since recruitment into the trial. The integer part is the number of days since recruitment.
- motor_UPDRS - Clinician's motor UPDRS score, linearly interpolated.
- ctotat_UPDRS - Clinician's total UPDRS score, linearly interpolated.
- Jitter(%), Jitter(Abs), Jitter:RAP, Jitter:PPQ5, Jitter:DDP - Several measures of variation in fundamental frequency.
- Shimmer, Shimmer(dB), Shimmer:APQ3, Shimmer:APQ5, Shimmer:APQ11, Shimmer:DDA - Several measures of variation in amplitude.
- NHR, HNR - Two measures of ratio of noise to tonal components in the voice.
- RPDE - A nonlinear dynamical complexity measure.
- DFA - Signal fractal scaling exponent.
- PPE - A nonlinear measure of fundamental frequency variation.

As noted, the objective is to predict the severity of the disease, where severity is defined based on the variable **total_UPDRS**: The disease is severe if **total_UPDRS** > **25**. This variable is created below.

4 Create the Binary Variable of Parkinson's Severity

```
data$severity <- data$total_UPDRS > 25
dim(data)
```

```
#> [1] 5875 23
```

```
summary(data$severity)
```

```
#>   Mode  FALSE   TRUE
#> logical    2188   3687
```

5 Normalization

The variables of the 16 voice measurements are normalized by means of the min-max transformation.

```
(biom_voice_mesures <- names(data)[7:22])
```

```
#> [1] "Jitter..."      "Jitter.Abs."      "Jitter.RAP"       "Jitter.PPQ5"
#> [5] "Jitter.DDP"       "Shimmer"          "Shimmer.dB."      "Shimmer.APQ3"
#> [9] "Shimmer.APQ5"     "Shimmer.APQ11"    "Shimmer.DDA"      "NHR"
#> [13] "HNR"              "RPDE"             "DFA"              "PPE"
```

```
mydata <- data[, biom_voice_mesures]
```

```
normalize <- function(x) {
  return((x- min(x))/(max(x)-min(x)))
}
```

```
mydata <- as.data.frame(lapply(mydata, normalize))
summary(mydata)
```

```
#>   Jitter...      Jitter.Abs.      Jitter.RAP      Jitter.PPQ5
#> Min.   :0.00000 Min.   :0.00000 Min.   :0.00000 Min.   :0.00000
#> 1st Qu.:0.02773 1st Qu.:0.04553 1st Qu.:0.02185 1st Qu.:0.02011
#> Median :0.04104 Median :0.07281 Median :0.03356 Median :0.02980
#> Mean   :0.05369 Mean   :0.09423 Mean   :0.04645 Mean   :0.04118
#> 3rd Qu.:0.06021 3rd Qu.:0.11523 3rd Qu.:0.05174 3rd Qu.:0.04383
#> Max.   :1.00000 Max.   :1.00000 Max.   :1.00000 Max.   :1.00000
#>   Jitter.DDP      Shimmer      Shimmer.dB.      Shimmer.APQ3
#> Min.   :0.00000 Min.   :0.00000 Min.   :0.00000 Min.   :0.00000
#> 1st Qu.:0.02185 1st Qu.:0.06047 1st Qu.:0.0716 1st Qu.:0.04762
#> Median :0.03361 Median :0.09207 Median :0.1091 Median :0.07507
#> Mean   :0.04650 Mean   :0.11664 Mean   :0.1369 Mean   :0.09652
#> 3rd Qu.:0.05179 3rd Qu.:0.13816 3rd Qu.:0.1629 3rd Qu.:0.11775
#> Max.   :1.00000 Max.   :1.00000 Max.   :1.00000 Max.   :1.00000
#>   Shimmer.APQ5      Shimmer.APQ11      Shimmer.DDA      NHR
#> Min.   :0.00000 Min.   :0.00000 Min.   :0.00000 Min.   :0.00000
#> 1st Qu.:0.05361 1st Qu.:0.04827 1st Qu.:0.04758 1st Qu.:0.01426
#> Median :0.08481 Median :0.07407 Median :0.07507 Median :0.02428
#> Mean   :0.11027 Mean   :0.09155 Mean   :0.09650 Mean   :0.04256
#> 3rd Qu.:0.13215 3rd Qu.:0.11073 3rd Qu.:0.11775 3rd Qu.:0.04168
#> Max.   :1.00000 Max.   :1.00000 Max.   :1.00000 Max.   :1.00000
#>   HNR      RPDE      DFA      PPE
#> Min.   :0.0000 Min.   :0.0000 Min.   :0.0000 Min.   :0.0000
#> 1st Qu.:0.4900 1st Qu.:0.3911 1st Qu.:0.2336 1st Qu.:0.1893
```

```
#> Median :0.5594    Median :0.4800    Median :0.3685    Median :0.2586
#> Mean   :0.5528    Mean    :0.4790    Mean    :0.3959    Mean    :0.2784
#> 3rd Qu.:0.6291    3rd Qu.:0.5681    3rd Qu.:0.5612    3rd Qu.:0.3417
#> Max.   :1.0000    Max.    :1.0000    Max.    :1.0000    Max.    :1.0000
```

6 Separation into Train and Test Data

I will use (pseudo-) random sampling to separate the data into a training and test set.

```
#set.seed(1)
ratio <- 0.7
sample.size <- floor(nrow(data) * ratio)
train.indices <- sample(1:nrow(data), size = sample.size)
train <- mydata[train.indices, ]
test <- mydata[-train.indices, ]

x_train <- data.matrix(train)
y_train <- as.numeric(data[train.indices,]$severity)
x_test <- data.matrix(test)
y_test <- as.numeric(data[-train.indices,]$severity)
```

7 Implementation of a Dense DNN

A dense deep neural network (DNN) for severity prediction is made. Using the `tfruns` package, we will determine which of the following architectures provides the best accuracy:

- Model 1: two hidden layers of 10 nodes each.
- Model 2: three hidden layers of 20, 10 and 5 nodes respectively.
- adding 40% dropout between hidden layers.
- adding l2 regularization on hidden layers.

7.1 Using One Output Node for Classification

Since we have one output node, we should use the *sigmoid* activation function in the output and the *binary_crossentropy* loss function. Moreover, I will be testing model 1 and model 2, both with and without dropout regularization and l2 regularization (in total 6 models).

```
# Using tfruns to test all the alternative we are given in the task description.
for (model in 1:6){
  training_run('neural_net.R', flags = c(model = model))
}
```

The code that I used to build the different neural networks is displayed below.

```
# Set hyperparameter flags.
FLAGS <- flags(
  flag_integer("model", 1)
)

# Defining the model and layers.
if (FLAGS$model == 1){
  model <- keras_model_sequential() %>%
    layer_dense(units = 10, activation = 'relu', input_shape = c(ncol(x_train))) %>%
    layer_dense(units = 10, activation = 'relu') %>%
    layer_dense(units = 1, activation = 'sigmoid')
```

```

} else if (FLAGS$model == 2){
  model <- keras_model_sequential() %>%
    layer_dense(units = 20, activation = 'relu', input_shape = c(ncol(x_train))) %>%
    layer_dense(units = 10, activation = 'relu') %>%
    layer_dense(units = 5, activation = 'relu') %>%
    layer_dense(units = 1, activation = 'sigmoid')
} else if (FLAGS$model == 3){
  # Model 1 with dropout regularization between the layers.
  model <- keras_model_sequential() %>%
    layer_dense(units = 10, activation = 'relu', input_shape = c(ncol(x_train))) %>%
    layer_dropout (rate = 0.4) %>%
    layer_dense(units = 10, activation = 'relu') %>%
    layer_dropout (rate = 0.4) %>%
    layer_dense(units = 1, activation = 'sigmoid')
} else if (FLAGS$model == 4){
  # Model 2 with dropout regularization between the layers.
  model <- keras_model_sequential() %>%
    layer_dense(units = 20, activation = 'relu', input_shape = c(ncol(x_train))) %>%
    layer_dropout (rate = 0.4) %>%
    layer_dense(units = 10, activation = 'relu') %>%
    layer_dropout (rate = 0.4) %>%
    layer_dense(units = 5, activation = 'relu') %>%
    layer_dropout (rate = 0.4) %>%
    layer_dense(units = 1, activation = 'sigmoid')
} else if (FLAGS$model == 5){
  # Model 1 with l2 regularization on the hidden layers.
  model <- keras_model_sequential() %>%
    layer_dense(units = 10, activation = 'relu', input_shape = c(ncol(x_train)), kernel_regularizer = r
    layer_dense(units = 10, activation = 'relu', kernel_regularizer = regularizer_l2(l = 0.01)) %>%
    layer_dense(units = 1, activation = 'sigmoid')
} else (FLAGS$model == 6){
  # Model 2 with l2 regularization on the hidden layers.
  model <- keras_model_sequential() %>%
    layer_dense(units = 20, activation = 'relu', input_shape = c(ncol(x_train)), kernel_regularizer = r
    layer_dense(units = 10, activation = 'relu', kernel_regularizer = regularizer_l2(l = 0.01)) %>%
    layer_dense(units = 5, activation = 'relu', kernel_regularizer = regularizer_l2(l = 0.01)) %>%
    layer_dense(units = 1, activation = 'sigmoid')
}

summary(model)

# compile (define loss and optimizer)
model %>% compile(loss = 'binary_crossentropy',
                 optimizer = optimizer_rmsprop(),
                 metrics = c('accuracy'))

# train (fit)
history <- model %>% fit(x_train, y_train, epochs = 100,
                       batch_size = 256, validation_split = 0.2)

# plot
plot(history)

```

```

# evaluate on training data.
score.train <- model %>% evaluate(x_train, y_train, verbose = 0)

cat('Train loss:', score.train[1], '\n')
cat('Train accuracy:', score.train[2], '\n')

# evaluate on testing data.
score.test <- model %>% evaluate(x_test, y_test, verbose = 0)

cat('Test loss:', score.test[1], '\n')
cat('Test accuracy:', score.test[2], '\n')

```

Comparison using the `tfruns` package can be done with the commands given in the code below

```

latest_run() # Show the latest trained model.
ls_runs() # Show all trained models.
ls_runs(metric_val_accuracy > 0.94, order = metric_val_accuracy) # Show a selection of trained models.
compare_runs() # Compare two runs (can be specified as an argument in the function).

```

```

ls_runs()

#> Data frame: 6 x 23
#>
#>   run_dir      eval_ metric_loss metric_accuracy
#> 1 runs/2022-03-11T17-42-13Z 0.6661904      0.6536      0.6397
#> 2 runs/2022-03-11T17-41-53Z 0.6422822      0.6357      0.6598
#> 3 runs/2022-03-11T17-41-32Z 0.6094625      0.6207      0.6397
#> 4 runs/2022-03-11T17-41-12Z 0.6067647      0.6146      0.6683
#> 5 runs/2022-03-11T17-40-51Z 0.5697196      0.5594      0.7188
#> 6 runs/2022-03-11T17-40-30Z 0.5923944      0.5714      0.7112
#>   metric_val_loss metric_val_accuracy
#> 1          0.6762          0.6002
#> 2          0.6519          0.6209
#> 3          0.6217          0.6002
#> 4          0.6125          0.6829
#> 5          0.5818          0.7145
#> 6          0.5824          0.6987
#> # ... with 17 more columns:
#> #   flag_model, epochs, epochs_completed, metrics, model, loss_function,
#> #   optimizer, learning_rate, script, start, end, completed, output,
#> #   source_code, context, type, NA.

```

According to the list of all runs, the two first models (without either of the regularization techniques) have the smallest loss and largest accuracy. Thus, comparing those two models reveals that model 2 has the largest test accuracy (even though all these models are terrible in practice).

```

compare_runs(ls_runs(metric_val_accuracy > 0.69, order = metric_val_accuracy))

```

8 Add Callbacks

In order to improve the selected model (Model 2) further, we add the following callbacks, before re-training the model and making predictions.

```

checkpoint.filepath <- "checkpoint.h5"
callbacks_list<-list(
  callback_early_stopping(

```

```

    monitor = "accuracy",
    patience=30
),
callback_model_checkpoint(
    filepath = checkpoint.filepath,
    monitor = "accuracy",
    save_best_only = T
),
callback_reduce_lr_on_plateau(
    monitor="accuracy",
    factor = 0.001,
    patience = 8
)
)

```

```
#> Loaded Tensorflow version 2.7.1
```

The first callback is used to stop training when accuracy stops improving. The patience is set to 30, which means that the training is stopped after 45 consecutive epochs where accuracy is not improved.

The second callback is used to save the best model found during training. The weights of this model can then be loaded into the environment afterwards, to ensure that we are using the best model found during training.

The last callback is used to reduce the learning rate when the accuracy has stopped improving. The patience is set to 8, which means that the learning rate will be decreased by a factor (set to 0.1) after 15 consecutive epochs of no improvement in accuracy.

```

model <- keras_model_sequential() %>%
  layer_dense(units = 20, activation = 'relu', input_shape = c(ncol(x_train))) %>%
  layer_dense(units = 10, activation = 'relu') %>%
  layer_dense(units = 5, activation = 'relu') %>%
  layer_dense(units = 1, activation = 'sigmoid')

summary(model)

```

```
#> Model: "sequential"
```

```

#> -----
#> Layer (type)                Output Shape          Param #
#> =====
#> dense_3 (Dense)             (None, 20)            340
#>
#> dense_2 (Dense)             (None, 10)            210
#>
#> dense_1 (Dense)             (None, 5)             55
#>
#> dense (Dense)               (None, 1)             6
#>
#> =====
#> Total params: 611
#> Trainable params: 611
#> Non-trainable params: 0
#> -----

```

```

# compile (define loss and optimizer)
model %>% compile(loss = 'binary_crossentropy',
                 optimizer = optimizer_rmsprop(),

```



```

        metrics = c('accuracy'))

# train (fit)
set.seed(1)
history <- model %>% fit(x_train, y_train, epochs = 200,
                        batch_size = 256, validation_split = 0.2,
                        callbacks = callbacks_list)

final.model <- load_model_hdf5(checkpoint.filepath)

```

9 Predictions

```

# Predictions for one output node
y_pred <- model %>% predict(x_test) %>% `>`(0.5) %>% k_cast("int32")
y_pred <- as.array(y_pred)
y_pred2 <- as.array(final.model %>% predict(x_test) %>% `>`(0.5) %>% k_cast("int32"))
all.equal(y_pred, y_pred2)

```

```
#> [1] "Mean relative difference: 3"
```

```
(tab <- table("Predictions" = y_pred, "Labels" = y_test))
```

```

#>           Labels
#> Predictions  0   1
#>           0 232 100
#>           1 432 999

```

```
# accuracy in predictions (as shown with the "evaluate" above).
```

```
(tab[1]+tab[4])/sum(tab)
```

```
#> [1] 0.6982416
```

```
# Better to do this with the package "caret".
```

```
confusionMatrix(factor(y_pred), factor(y_test), positive = "1")
```

```
#> Confusion Matrix and Statistics
```

```

#>
#>           Reference
#> Prediction  0   1
#>           0 232 100
#>           1 432 999
#>
#>           Accuracy : 0.6982
#>           95% CI : (0.6762, 0.7196)
#>    No Information Rate : 0.6234
#>    P-Value [Acc > NIR] : 2.646e-11
#>
#>           Kappa : 0.2868
#>
#> Mcnemar's Test P-Value : < 2.2e-16
#>
#>           Sensitivity : 0.9090
#>           Specificity : 0.3494
#>    Pos Pred Value : 0.6981
#>    Neg Pred Value : 0.6988

```

```

#>           Prevalence : 0.6234
#>       Detection Rate : 0.5666
#>   Detection Prevalence : 0.8117
#>       Balanced Accuracy : 0.6292
#>
#>       'Positive' Class : 1
#>
confusionMatrix(factor(y_pred2), factor(y_test), positive = "1")

#> Confusion Matrix and Statistics
#>
#>           Reference
#> Prediction    0    1
#>           0 225 102
#>           1 439 997
#>
#>           Accuracy : 0.6931
#>           95% CI : (0.671, 0.7146)
#>   No Information Rate : 0.6234
#>   P-Value [Acc > NIR] : 5.119e-10
#>
#>           Kappa : 0.2735
#>
#>   McNemar's Test P-Value : < 2.2e-16
#>
#>           Sensitivity : 0.9072
#>           Specificity : 0.3389
#>       Pos Pred Value : 0.6943
#>       Neg Pred Value : 0.6881
#>           Prevalence : 0.6234
#>       Detection Rate : 0.5655
#>   Detection Prevalence : 0.8145
#>       Balanced Accuracy : 0.6230
#>
#>       'Positive' Class : 1
#>

```

The accuracy is still not good for this model, even though it perhaps is the best model among the ones we tested.